

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from IPython.display import display
from pylab import rcParams
from datetime import datetime, timedelta
from pandas.tseries.offsets import BDay
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.stattools import pacf
from statsmodels.tsa.stattools import acf
from statsmodels.graphics.tsaplots import plot_pacf
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.gofplots import qqplot
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.api import ExponentialSmoothing, Holt
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
import itertools
```

```
In [2]: plt.rcParams['figure.figsize'] = 12,8
```

```
In [3]: df = pd.read_csv('retail_sales.csv')
df.head()
```

```
Out [3]:
```

	ds	y
0	1992-01-01	146376
1	1992-02-01	147079
2	1992-03-01	159336
3	1992-04-01	163669
4	1992-05-01	170068

```
In [4]: df.head(),df.tail()
```

```
Out [4]: (
   ds      y
0 1992-01-01 146376
1 1992-02-01 147079
2 1992-03-01 159336
3 1992-04-01 163669
4 1992-05-01 170068,
   ds      y
288 2016-01-01 400928
289 2016-02-01 413554
290 2016-03-01 460093
291 2016-04-01 450935
292 2016-05-01 471421)
```

```
In [5]: date = pd.date_range(start='1992/01/01',end='2016/05/01',freq=BDay())
date
```

```
Out [5]: DatetimeIndex(['1992-01-01', '1992-01-02', '1992-01-03', '1992-01-06',
                        '1992-01-07', '1992-01-08', '1992-01-09', '1992-01-10',
                        '1992-01-13', '1992-01-14',
                        ...,
                        '2016-04-18', '2016-04-19', '2016-04-20', '2016-04-21',
                        '2016-04-22', '2016-04-25', '2016-04-26', '2016-04-27',
                        '2016-04-28', '2016-04-29'],
                        dtype='datetime64[ns]', length=6348, freq='B')
```


```
In [6]: df['time_stamp'] = pd.to_datetime(df['ds'],format='%Y-%m-%d')
df = df.set_index('time_stamp')
df = df.drop('ds',axis=1)
df.head()
```

Out [6]:

y

time_stamp	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
In [7]: df.plot()
plt.show()
```



```
In [8]: df.isnull().sum()
```

Out [8]: y 0
dtype: int64

```
In [9]: df.describe()
```

Out [9]:

y


count	293.000000
mean	308971.310580
std	86084.323897
min	146376.000000
25%	234503.000000
50%	314051.000000

y


75% 375795.000000

max 518253.000000

```
In [10]: decompose_add = seasonal_decompose(df, model='add')
decompose_add.plot()
plt.show()
```



```
In [11]: decompose_mul = seasonal_decompose(df, model='mul')
decompose_mul.plot()
plt.show()
```



```
In [12]: observation_df = df.values
result_df = adfuller(observation_df)

print(f"ADF test statistics: {result_df[0]}")
print(f"p_value: {result_df[1]}")
print(f"Used lags: {result_df[2]}")
print(f"Number of observations: {result_df[3]}")
for key,value in result_df[4].items():
    print(f"{key}: {value}")

print(f"IC best: {result_df[5]}")
```

```
ADF test statistics: -0.6443250505938629
p_value: 0.8606486181958646
Used lags: 15
Number of observations: 277
1%: -3.4541800885158525
5%: -2.872031361137725
10%: -2.5723603999791473
IC best: 5671.009371066914
```


```
In [13]: df_diff = df.diff(periods=1).dropna()
observation_diff = df_diff.values
result_diff = adfuller(observation_diff)

print(f"ADF test statistics: {result_diff[0]}")
print(f"p_value: {result_diff[1]}")
print(f"used lags: {result_diff[2]}")
print(f"Number of observations: {result_diff[3]}")
for key,values in result_diff[4].items():
    print(f"{key}: {values}")
print(f"IC Best: {result_diff[5]}")
```


```
ADF test statistics: -3.437937154315123
p_value: 0.009731610960261502
used lags: 14
Number of observations: 277
1%: -3.4541800885158525
5%: -2.872031361137725
```

10%: -2.5723603999791473
IC Best: 5649.985617927769

```
In [14]: plot_acf(df, lags=30)  
plt.show()
```



```
In [15]: plot_pacf(df, lags=30)  
plt.show()
```



```
In [16]: df.head(),df.tail()
```

```
Out [16]: (
      time_stamp      y
1992-01-01  146376
1992-02-01  147079
1992-03-01  159336
1992-04-01  163669
1992-05-01  170068,
      time_stamp      y
2016-01-01  400928
2016-02-01  413554
2016-03-01  460093
2016-04-01  450935
2016-05-01  471421)
```

```
In [17]: train_set = datetime(2013,11,1)
test_set = datetime(2016,5,1)
train = df[:train_set]
test = df[train_set+timedelta(days=1):test_set]
```

```
In [18]: train.shape,test.shape
```

```
Out [18]: ((263, 1), (30, 1))
```

```
In [19]: train.isnull().sum()
```

```
Out [19]: y      0
dtype: int64
```

```
In [20]: p = q = range(0,4)
d = range(0,2)
pdq = itertools.product(p,d,q)
arima_model = []
for i in pdq:
    model = ARIMA(train,order=i)
    result_arima = model.fit()
    arima_model.append((i,result_arima.aic))
```

```
In [21]: frame = pd.DataFrame(arima_model,columns=['param', 'aic'])
```

```
In [22]: frame.sort_values(by='aic',ascending=True).head()
```

```
Out [22]:
```

	param	aic
27	(3, 0, 3)	6000.167826
19	(2, 0, 3)	6012.776444
15	(1, 1, 3)	6014.462657
23	(2, 1, 3)	6015.041477
11	(1, 0, 3)	6015.902494

```
In [23]: model_arima = ARIMA(train,order=(3,0,3))
result_arima_1 = model_arima.fit()
print(result_arima_1.summary())
```

```

              SARIMAX Results
=====
Dep. Variable:              y      No. Observations:              263
Model:              ARIMA(3, 0, 3)  Log Likelihood              -2992.084
Date:              Wed, 19 Feb 2025  AIC              6000.168
Time:              22:07:04         BIC              6028.745
Sample:              01-01-1992     HQIC              6011.652
              - 11-01-2013
Covariance Type:              opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
=====
```

```

const      2.94e+05  1.41e-09  2.08e+14  0.000  2.94e+05  2.94e+05
ar.L1      -0.6535  0.050  -12.999  0.000  -0.752  -0.555
ar.L2       0.7872  0.022  35.956  0.000  0.744  0.830
ar.L3       0.8660  0.048  17.986  0.000  0.772  0.960
ma.L1       1.0641  0.079  13.476  0.000  0.909  1.219
ma.L2      -0.4956  0.085  -5.842  0.000  -0.662  -0.329
ma.L3      -0.7968  0.072  -11.056  0.000  -0.938  -0.656
sigma2     4.862e+08  1.36e-10  3.58e+18  0.000  4.86e+08  4.86e+08
=====
Ljung-Box (L1) (Q):      3.03  Jarque-Bera (JB):      12.21
Prob(Q):                0.08  Prob(JB):              0.00
Heteroskedasticity (H):  2.11  Skew:              -0.10
Prob(H) (two-sided):    0.00  Kurtosis:         4.04
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 4.51e+33. Standard errors may be unstable.

```
In [24]: from sklearn.metrics import root_mean_squared_error
```

```
In [25]: class Evaluate:
        def Rmse(y_true,y_pred):
            return root_mean_squared_error(y_true,y_pred)

        def Mape(y_true,y_pred):
            return np.mean((np.abs(y_true-y_pred))/y_true)*100
```

```
In [26]: arima_pred = result_arima_1.forecast(len(test))
```


```
In [27]: Evaluate.Rmse(test['y'],arima_pred)
```

Out [27]: 38866.408528879765

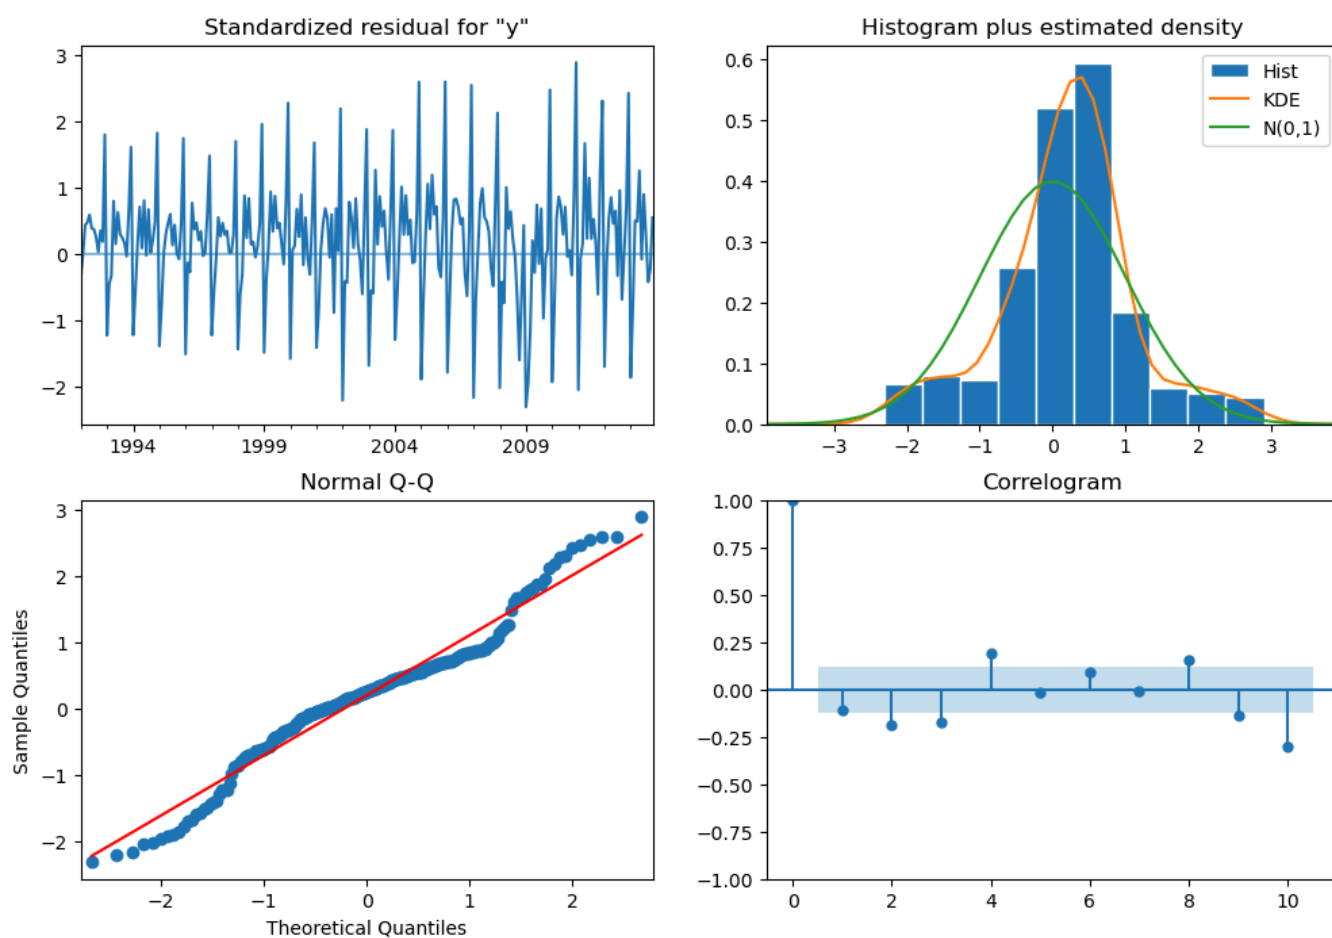
```
In [28]: Evaluate.Mape(test['y'],arima_pred)
```

Out [28]: 7.309302838308487

```
In [29]: plt.plot(train)
        plt.plot(test)
        plt.plot(arima_pred)
        plt.show()
```



```
In [30]: result_arima_1.plot_diagnostics()
plt.show()
```



```
In [31]: model_exp = ExponentialSmoothing(train, trend='additive', seasonal='additive', initialization='method_of_moments')
result_exp = model_exp.fit()
print(result_exp.summary())
```

```
=====
ExponentialSmoothing Model Results
=====
Dep. Variable:          y    No. Observations:      263
```



```


Model: ExponentialSmoothing SSE 10327972264.770
Optimized: True AIC 4630.810
Trend: Additive BIC 4687.964
Seasonal: Additive AICC 4633.613
Seasonal Periods: 12 Date: Wed, 19 Feb 2025
Box-Cox: False Time: 22:07:05
Box-Cox Coeff.: None

```

	coeff	code	optimized
smoothing_level	0.5000000	alpha	True
smoothing_trend	0.0001	beta	True
smoothing_seasonal	0.3250000	gamma	True
initial_level	1.6654e+05	l.0	True
initial_trend	972.51187	b.0	True
initial_seasons.0	-24263.224	s.0	True
initial_seasons.1	-24717.464	s.1	True
initial_seasons.2	-1087.3281	s.2	True
initial_seasons.3	-2654.0052	s.3	True
initial_seasons.4	6487.0990	s.4	True
initial_seasons.5	3994.2552	s.5	True
initial_seasons.6	739.52604	s.6	True
initial_seasons.7	4860.1094	s.7	True
initial_seasons.8	-4216.2552	s.8	True
initial_seasons.9	-193.24479	s.9	True
initial_seasons.10	2540.8698	s.10	True
initial_seasons.11	38509.661	s.11	True

```
In [32]: pred_exp = result_exp.forecast(len(test))
```

```
In [33]: plt.plot(train)
plt.plot(test)
plt.plot(pred_exp)
plt.show()
```



```
In [34]: Evaluate.Rmse(test['y'], pred_exp)
```


```
Out [34]: 8054.220053734076
```

```
In [35]: Evaluate.Mape(test['y'], pred_exp)
```

```
Out [35]: 1.456057930754873
```

```
In [36]: residuals = test['y'].values - pred_exp.values
```

```
In [37]: qqplot(residuals,line='s')
plt.show()
```



```
In [38]: model_exp_tune = ExponentialSmoothing(train['y'],trend='add',seasonal='mul',seasonal_perio
result_exp_tune = model_exp_tune.fit()
print(result_exp_tune.summary())
```

```


ExponentialSmoothing Model Results
=====
Dep. Variable:          y      No. Observations:          263
Model:      ExponentialSmoothing  SSE          9936302247.368
Optimized:      True      AIC          4620.642
Trend:      Additive      BIC          4677.796
Seasonal:      Multiplicative  AICC          4623.445
Seasonal Periods:      12      Date:      Wed, 19 Feb 2025
Box-Cox:      False      Time:      22:07:06
Box-Cox Coeff.:      None
=====

      coeff      code      optimized
-----
smoothing_level      0.5000000      alpha      True
smoothing_trend      0.0001      beta      True
smoothing_seasonal      0.2750000      gamma      True
initial_level      1.6654e+05      l.0      True
initial_trend      972.51187      b.0      True
initial_seasons.0      0.8743406      s.0      True
initial_seasons.1      0.8715923      s.1      True
initial_seasons.2      0.9938863      s.2      True
initial_seasons.3      0.9869043      s.3      True
initial_seasons.4      1.0320753      s.4      True
initial_seasons.5      1.0201988      s.5      True
initial_seasons.6      1.0045706      s.6      True
initial_seasons.7      1.0248818      s.7      True
initial_seasons.8      0.9773991      s.8      True
initial_seasons.9      0.9998626      s.9      True
initial_seasons.10      1.0129029      s.10     True
initial_seasons.11      1.2013853      s.11     True
=====
```

```
In [39]: final_exp = ExponentialSmoothing(df['y'],trend='add',seasonal='add',seasonal_periods=12)
result_final = final_exp.fit(smoothing_level=0.5,
                             smoothing_trend = 0.0001,
                             smoothing_seasonal = 0.2750000 ,
                             optimized = False)
```

```
In [40]: final_pred = result_final.forecast(30)
```

```
In [41]: plt.plot(train)
plt.plot(test)
plt.plot(final_pred)
plt.show()
```



```
In [42]: new = pd.DataFrame(final_pred, columns=['pred'])
new.head()
```

```
Out [42]:
```

	pred
2016-06-01	454690.460806
2016-07-01	460684.693340
2016-08-01	466657.798203
2016-09-01	438619.823755
2016-10-01	452555.354617

```
In [ ]:
```